

## Diamond Price OLS / SGD

### **Installs:**

```
pip install pandas
pip install scikit-learn
pip install matplotlib
pip install seaborn
```

### **Dataset:**

<https://www.kaggle.com/datasets/shivam2503/diamonds>

I used a diamond price dataset (53.9k entries) which is a collection of diamond dimensions and carats, with their prices. All features: carat, depth-percentage, width-percentage, length, width, depth, cut, color, and clarity. But, I only manually extracted six continuous features (carat, depth-percentage, width-percentage, length, width, depth) from this dataset. The continuous output was the diamond's price. I loaded the dataset using the pandas library, in which I use the function `read_csv()`.

**Training/Test Data Splits:** I used the `sklearn.model_selection` function `train_test_split()` to use a random 20% of the dataset for testing, and the remaining for training.

**OLS Data Preparation:** I appended an `x0` column of 1's to be the first row of the dataset. This allows the OLS matrix operations to generate a bias (intercept) term in the set of parameters.

**OLS Model Training:** I first converted the pandas training/test & input/output datasets to numpy for matrix operations. I then transposed the training input using `np.matrix.transpose()`, and multiplied it to the original using `np.matmul()`, I then got the inverse of the product using `np.linalg.inv()`. I then calculated the product of the prior operations with the training input transposed and the training output transposed to get the optimal parameters.

```
OLS Linear Equation:
price = 20976.563736020486*x0 + 10683.18425965889*carat + -204.09955335264056*depth% + -104.26736871735861*width% + -1286.7995742830
012*Length + 37.63252068422567*Width + 53.3948642392736*Depth
```

**OLS Model Predictions:** I used `np.matmul()` to get the product of the OLS parameters and the training and test inputs individually to get an array of training and testing output predictions respectively.

**OLS Model Evaluation:** I used the `sklearn.metrics` function `mean_squared_error()` to pass the training & test data with its real-valued predictions to get the mean squared error: **Training MSE** = 2240338, **Test MSE** = 2242178. I also used the `sklearn.metrics` function `mean_absolute_error()` to get the mean absolute error: **Training MAE** = 891, **Test MAE** = 888. I also used the `sklearn.metrics` function `r2_score()` to get the coefficient of determination: **Training R2** = 0.8593, **Test R2** = 0.8590.

**SGD Model Training:** I first created a pipeline (sklearn.pipeline) with make\_pipeline(), StandardScaler(), and SGDRegressor() to scale and organize my data. In SGDRegressor I passed the default parameters max\_iter=1000 to clarify the number of passes through the training set, and tol=1e-3 for the stopping criterion. This is a gradient descent model that dynamically changes the learning rate as we reach an optimal solution.

**SGD Model Predictions:** Using the SGDRegressor model from before and the predict() function with the training and test inputs individually to get an array of training and testing output predictions respectively.

**SGD Model Evaluation:** I used the sklearn.metrics function mean\_squared\_error() to pass the training & test data with its real-valued predictions to get the mean squared error: **Training MSE** = 2242415, **Test MSE** = 2245159.

I also used the sklearn.metrics function mean\_absolute\_error() to get the mean absolute error: **Training MAE** = 892, **Test MAE** = 889.

I also used the sklearn.metrics function r2\_score() to get the coefficient of determination: **Training R2** = 0.8591, **Test R2** = 0.8588.

**NEXT PAGE**

### Diamond Price Correlation Heat Map:



I used the `corr()` function to evaluate the dataset, and `matplotlib.pyplot` / `seaborn` to visualize the correlation matrix. As you can see there is definitely a correlation between diamond price and length, width, height, and carat. You can also see that width% or fatter diamonds sold for more.